

TURNING PUBLIC REVIEWS INTO PRODUCT SIGNAL

METHOD PRINCIPLE

Preserve source facts, check the meaning of every classification, challenge each theme, and separate daily monitoring from human-approved executive conclusions.

Seven-phase evidence path

	Phase	Decision	Gate or boundary
01	Define the question and source boundaries	Use all unique written reviews available from the four public sources, while treating storefront feedback as directional evidence.	Apple Driver exposes a capped feed and Google listing totals do not always match exhausted pagination. Unique extracted review IDs are the count of record.
02	Build a trustworthy extraction and storage foundation	Separate raw evidence, normalized facts, and analytical labels instead of overwriting source data.	A failed pull cannot replace the last good state, and raw reviewer data remains private.
03	Profile the data and choose comparable windows	Use September 1, 2025 through August 31, 2026 for the approved analysis and use separate earlier periods for recent-trend comparisons.	Google Driver dominates the natural sample, so rates use the relevant product or platform denominator rather than the combined total.
04	Discover, align, and calibrate the taxonomy	Use a structured taxonomy for product area, issue, customer signal, failure mode, and potential consequence.	An out-of-window 2022 review was removed before coding so calibration matched the approved population.
05	Reopen and correct the semantic classification gate	Rebuild the reliable analysis base instead of accepting technically valid but semantically weak labels.	The manual feedback-strength label did not meet the reproducibility threshold, so it was narrowed to a conservative confirmation flag rather than weakening the gate.
06	Synthesize, challenge, and approve findings	Publish only themes that survive the five-part evidence review and receive human approval.	Public reviews support where to investigate and improve. They do not establish prevalence across all customers or prove root cause.
07	Lift the analysis into a shareable, repeatable system	Run two clocks: human-approved analytical releases for executive content and automated monitoring for newly observed signals.	New reviews can update the monitor automatically, but themes, recommendations, and executive claims change only through review and a new approved release.

Evidence review: what a theme must pass

01 DENOMINATOR Correct Fleet or Driver base	02 TIME Non-overlapping periods; no review on both sides	03 PLATFORM Like-for-like Android and iOS comparison	04 CHALLENGE Counterevidence and alternative explanations	05 APPROVAL Human sign-off and an explicit claim boundary
--	---	---	--	--

Fleet example: stability is 42.7% Android versus 20.0% iOS across all Fleet reviews. Among reviews that already report a problem, it is 49.3% versus 45.0%. Android has a more negative review population overall, but the problem mix is much closer. Report concentrated negative Android feedback, not proof of a different Android product problem.

Four judgment moments

Judgment	Action taken
Do not force the source totals to match	Count unique review IDs actually extracted and disclose endpoint limits.
Correct the sample before coding	Remove it and rebuild the sample before taxonomy work continued.
Reopen a gate that passed technically	Pause downstream analysis, correct carried records, blind-review the rest, and rerun consistency testing.
Narrow the Fleet Android interpretation	Compare the theme among reviews that already signal a problem. The gap narrowed from 42.7% versus 20.0% of all reviews to 49.3% versus 45.0% of problem-signal reviews.

Operating model, limits, and next week

Automated	Human reviewed	Confidence, AI + access	With another week
Collect, normalize, and compare; run integrity, privacy, and publication gates; refresh the monitor and preserve the last good state.	Approve taxonomy and themes; challenge claims and recommendations; approve executive releases.	Confident in nine recurring patterns; not in prevalence, root cause, or release causality. AI supported analysis and drafting; human checks governed claims. No owner-console or internal data. V3: 924 base +5 eligible -2 aged out = 927.	Join telemetry and support; validate owners and thresholds; repair response-time anomalies; activate the private schedule.